


Approximate Near Neighbor Search: Locality Sensitive Hashing

Inge Li Gørtz


Nearest Neighbor

- **Nearest Neighbor.** Given a set of points P in a metric space, build a data structure which given a query point x returns the point in P closest to x .
- **Metric.** Distance function d is a metric:
 1. $d(x,y) \geq 0$
 2. $d(x,y) = 0$ if and only if $x = y$
 3. $d(x,y) = d(y,x)$
 4. $d(x,y) \leq d(x,z) + d(z,y)$
- **Warmup.** 1D: Real line



Approximate Near Neighbors

- **ApproximateNearNeighbor(x):** Return a point y such that $d(x, y) \leq c \cdot \min_{z \in P} d(x, z)$
- **c-Approximate r-Near Neighbor:** Given a point x if there exists a point z in P such that $d(x, z) \leq r$ then return a point y such that $d(x, y) \leq c \cdot r$. If no such point z exists return Fail.
- Randomised version: Return such an y with probability δ .




Locality Sensitive Hashing

- **Locality sensitive hashing.** A family of hash functions H is (r, cr, p_1, p_2) -sensitive with $p_1 > p_2$ and $c > 1$ if:

- $d(x, y) \leq r \Rightarrow P[h(x) = h(y)] \geq p_1$ (close points)
- $d(x, y) \geq cr \Rightarrow P[h(x) = h(y)] \leq p_2$ (distant points)

for h chosen randomly from H .



Hamming Distance

- P set of n bit strings each of length d.
- **Hamming distance.** the number of bits where x and y differ:

$$d(x, y) = |\{i : x_i \neq y_i\}|$$

- Example.

$$\begin{array}{r} x = \boxed{1} \boxed{0} 1 0 0 1 \boxed{0} 0 \\ y = \boxed{0} \boxed{1} 1 0 0 1 \boxed{1} 0 \end{array} \quad \text{Hamming distance} = 3$$

- **Hash function.** Choose $i \in \{1, \dots, d\}$ uniformly at random and set $h(x) = x_i$.
- What is the probability that $h(x) = h(y)$?
 - $d(x, y) \leq r \Rightarrow P[h(x) = h(y)] \geq 1 - r/d$
 - $d(x, y) \geq cr \Rightarrow P[h(x) = h(y)] \leq 1 - cr/d$

LSH with Hamming Distance: Solution 1

- Pick random index i uniformly at random. Let $h(x) = x_i$.
- Bucket: Strings with same hash value $h(x)$.
- **Insert(x)**: Insert x in the list $A[h(x)]$
- **NearNeighbour(x)**: Compute Hamming distance from x to all bitstrings in $A[h(x)]$ until find one that is at most cr away. If no such string found return FAIL.

$$h(x) = x_3$$

$$a = 0011101$$

$$h(a) = 1$$

$$d = 0110011$$

$$h(d) = 1$$

$$b = 0101001$$

$$h(b) = 0$$

$$e = 1011101$$

$$h(e) = 1$$




$$c = 0010010$$

$$h(c) = 1$$

$$f = 1101101$$

$$h(f) = 0$$

Query time: $O(nd)$.



LSH with Hamming Distance: Solution 2

- Pick k random indexes uniformly and independently at random with replacement:

- $g(x) = x_{i_1}x_{i_2}\cdots x_{i_k}$

- Example. $k = 3$. $g(x) = x_2x_3x_6$

$$\begin{array}{r} x = \quad 1 \quad \boxed{0} \quad \boxed{1} \quad 0 \quad 0 \quad \boxed{1} \quad 0 \quad 0 \\ y = \quad 0 \quad \boxed{1} \quad \boxed{1} \quad 0 \quad 0 \quad \boxed{1} \quad 1 \quad 0 \end{array}$$


$$\begin{aligned} g(x) &= 011 \\ g(y) &= 111 \end{aligned}$$

y hashes to same value as x with probability **at most** p_2^k

- Probability that $g(x) = g(y)$?

- $d(x, y) \leq r \Rightarrow P[g(x) = g(y)] \geq (1 - r/d)^k$
- $d(x, y) \geq cr \Rightarrow P[g(x) = g(y)] \leq (1 - cr/d)^k$

z hashes to same value as x with probability **at least** p_1^k



no guarantees

LSH with Hamming Distance: Solution 2





- Pick k random indexes uniformly and independently at random with replacement:
 - $g(x) = x_{i_1}x_{i_2}\cdots x_{i_k}$
- Bucket: Strings with same hash value $g(x)$.

$$g(x) = x_2x_4x_7$$

$$a = 0011101 \quad g(a) = 011 \quad d = 0110011 \quad g(d) = 101$$

$$b = 0101001 \quad g(b) = 111 \quad e = 1011101 \quad g(e) = 011$$

$$c = 0010010 \quad g(c) = 000 \quad f = 1101101 \quad g(f) = 111$$



LSH with Hamming Distance: Solution 2

- Pick k random indexes uniformly and independently at random with replacement:

$$\bullet \quad g(x) = x_{i_1}x_{i_2}\cdots x_{i_k}$$

$$h_T(011_2) = 1$$

- Bucket: Strings with same hash value $g(x)$.

$$h_T(111_2) = 6$$

- Save buckets in a hash table T with hash function h_T .

$$h_T(000_2) = 9$$






$$h_T(101_2) = 1$$

$$g(x) = x_2x_4x_7$$

$$a = 0011101 \quad g(a) = 011 \quad d = 0110011 \quad g(d) = 101$$

$$b = 0101001 \quad g(b) = 111 \quad e = 1011101 \quad g(e) = 011$$

$$c = 0010010 \quad g(c) = 000 \quad f = 1101101 \quad g(f) = 111$$



LSH with Hamming Distance: Solution 2

- Pick k random indexes uniformly and independently at random with replacement:

$$\cdot g(x) = x_{i_1}x_{i_2}\cdots x_{i_k}$$

$$h_T(011_2) = 1$$

- Bucket: Strings with same hash value $g(x)$.

$$h_T(111_2) = 6$$

- Save buckets in a hash table T with hash function h_T .

$$h_T(000_2) = 9$$

- Insert(x):** Insert x in the list of $g(x)$ in T .

$$h_T(101_2) = 1$$






- NearNeighbour(x):** Compute Hamming distance from x to all bitstrings in $g(x)$ until find one that is at most cr away. If no such string found return FAIL.

$$g(x) = x_2x_4x_7$$


$$a = 0011101 \quad g(a) = 011 \quad d = 0110011 \quad g(d) = 101$$

$$b = 0101001 \quad g(b) = 111 \quad e = 1011101 \quad g(e) = 011$$

$$c = 0010010 \quad g(c) = 000 \quad f = 1101101 \quad g(f) = 111$$




LSH with Hamming Distance: Solution 2




- What happens when we increase k ?
 - Far away strings:

LSH with Hamming Distance: Solution 2




- What happens when we increase k ?
 - Far away strings: Probability that a far away string hashes to the same bucket as x decrease.

LSH with Hamming Distance: Solution 2




- What happens when we increase k ?
 - Far away strings: Probability that a far away string hashes to the same bucket as x decrease.
 - Close strings:

LSH with Hamming Distance: Solution 2




- What happens when we increase k ?
 - Far away strings: Probability that a far away string hashes to the same bucket as x decrease.
 - Close strings: Probability that a close string hashes to the same as x decrease.

LSH with Hamming Distance: Solution 2



- Expected number of far away strings that hash to same bucket as x :
 - $F = \{y : d(x, y) > cr\}$.
 - For $y \in F$ we want $P[g(y) = g(x)] \leq 1/n$:
 - Set $k = \lg n / \lg(1/p_2)$
 - $X_y = \begin{cases} 1 & y \text{ collides with } x \\ 0 & \text{otherwise} \end{cases}$
 - #far away strings colliding with x : $X = \sum_{y \in F} X_y$
 - $E[X] = \sum_{y \in F} E[X_y] = \sum_{y \in F} 1/n \leq 1$.
 - Markov: $P[X > 6] < E[X]/6 \leq 1/6$.



LSH with Hamming Distance: Solution 2



- What happens when we increase k ?
 - Probability that a far away string hashes to the same bucket as x decrease.
 - $k = \lg n / \lg(1/p_2)$ \Rightarrow with probability $\geq 5/6$ at most 6 far away strings hash to x 's bucket.
 - Probability that a close string hashes to the same as x decrease. 😞

LSH with Hamming Distance: Solution 3 (Amplification)

- Construct L hash tables T_j . Each table T_j has its own independently chosen **hash function** h_j and its own independently chosen **locality sensitive hash function** g_j .
- **Insert(x):** For all $1 \leq j \leq L$ insert x in the list of $g_j(x)$ in T_j .
- **Query(x):** For all $1 \leq j \leq L$ check each element in bucket $g_j(x)$ in T_j . Return the one closest to x if it is at most cr away. Otherwise return FAIL.



LSH with Hamming Distance

Let $k = \frac{\lg n}{\lg(1/p_2)}$, $\rho = \frac{\lg(1/p_1)}{\lg(1/p_2)}$, and $L = \lceil 2n^\rho \rceil$, where $p_1 = 1 - r/d$ and $p_2 = 1 - cr/d$.

- **Claim 1.** If there exists a string z^* in P with $d(x, z^*) \leq r$ then with probability at least $5/6$ we will return some z in P for which $d(x, z) \leq r$.

- Probability that z^* collides with x :

$$\cdot P[\exists i : g_i(x) = g_i(z^*)] = 1 - P[g_i(x) \neq g_i(z^*) \text{ for all } i]$$


$$= 1 - \prod_{i=1}^L P[g_i(x) \neq g_i(z^*)]$$

$$= 1 - \prod_{i=1}^L (1 - P[g_i(x) = g_i(z^*)])$$

$$\geq 1 - \prod_{i=1}^L (1 - p_1^k) = 1 - (1 - p_1^k)^L \geq 1 - e^{-Lp_1^k}$$


$$\geq 1 - \frac{1}{e^2} \geq 1 - 1/6 = 5/6$$

z^* hashes to same value as x with probability **at least** p_1^k



LSH with Hamming Distance

- Fast query time.
 - Check at most $6L + 1$ strings and return FAIL if no close string found.
 - Otherwise return closest string found.
- Theorem. If there exists a string z^* in P with $d(x, z^*) \leq r$ then with probability at least $2/3$ we will return some y in P for which $d(x, y) \leq cr$.
- Proof idea.
 - Show that with probability at least $5/6$ there are at most $6L$ far away strings that collides with x .
 - Already showed the probability that z^* is in the same bucket as x in at least one of the L hash tables is at least $5/6$.



LSH with Hamming Distance

- Insert time $O(kL)$.
- Expected query time $O(L(k + d))$.
 - $O(L)$ checks.
 - Each check takes $O(k + d)$ time.

Locality Sensitive Hashing

- **Locality sensitive hash function.** A family of hash functions \mathcal{H} is (r, cr, p_1, p_2) -sensitive with $p_1 > p_2$ and $c > 1$ if:
 - $d(x, y) \leq r \Rightarrow P[h(x) = h(y)] \geq p_1$ (close points)
 - $d(x, y) \geq cr \Rightarrow P[h(x) = h(y)] \leq p_2$ (distant points)
- **Amplification.**
 - Choose L hash functions $g_j(x) = h_{1,j}(x) \cdot h_{2,j}(x) \cdots h_{k,j}(x)$, where $h_{i,j}$ is chosen independently and uniformly at random from \mathcal{H} .
- **Locality sensitive hashing scheme.**
 - Construct L hash tables T_j .
 - **Insert(x):** For all $1 \leq j \leq L$ insert x in the list of $g_j(x)$ in T_j .
 - **Query(x):** For all $1 \leq j \leq L$ check each element in bucket $g_j(x)$ in T_j . Return the one closest to x . Check at most $6L + 1$ elements. If no element found at distance less than $c \cdot r$ from x return FAIL.

Jaccard distance and Min Hash

- **Jaccard distance.** Jaccard similarity: $\text{Jsim}(A, B) = \frac{|A \cap B|}{|A \cup B|}$
 - Jaccard distance: $1 - \text{Jsim}(A, B)$.
 - Hash function: *Min Hash*. (exercise)


Exercises

Angular Distance and Sim Hash

- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi.$
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$

Angular Distance and Sim Hash


- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi$.
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$



a				
b				
c				
d				
e				

Angular Distance and Sim Hash


- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi$.
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$



a	blue				
b					
c					
d	red				
e	red				

Angular Distance and Sim Hash


- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi$.
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$



a	Blue	Blue	White	White	White
b	Blue	Red	White	White	White
c	Blue	Red	White	White	White
d	Red	Red	Blue	White	White
e	Red	Red	Red	White	White

Angular Distance and Sim Hash


- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi$.
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$



a	blue	blue	blue		
b	blue	red	blue		
c	blue	red	blue		
d	red	blue	blue		
e	red	red	red		

Angular Distance and Sim Hash


- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi$.
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$



a	Blue	Blue	Blue	Blue	White
b	Blue	Red	Blue	Blue	White
c	Blue	Red	Blue	Blue	White
d	Red	Blue	Blue	Red	White
e	Red	Red	Red	Blue	White

Angular Distance and Sim Hash


- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi$.
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$



a	Blue	Blue	Blue	Blue	Blue
b	Blue	Red	Blue	Blue	Blue
c	Blue	Red	Blue	Blue	Blue
d	Red	Blue	Blue	Red	Red
e	Red	Red	Red	Blue	Red

Angular Distance and Sim Hash

- Collection of vectors.
- Distance between two vectors is the angular distance between them
 $\text{dist}(u, v) = \angle(u, v)/\pi.$
 - Assume u and v are unit vectors. Then $u \cdot v = \cos(\angle(u, v))$
- Hash function: Sim Hash.
 - Random projection: Take a random vector r and set $h_r(u) = \text{sign}(r \cdot u)$



a	Blue	Blue	Blue	Blue	Blue
b	Blue	Red	Blue	Blue	Blue
c	Blue	Red	Blue	Blue	Blue
d	Red	Blue	Blue	Red	Red
e	Red	Red	Red	Blue	Red

- Can show that $P[h(u) = h(v)] = 1 - \angle(u, v)/\pi.$