

Dynamic Connectivity

Christian Wulff-Nilsen

Algorithmic Techniques for Modern Data Models

DTU

November 28, 2025


Problem Definition

- Obtain an efficient data structure supporting the following operations in a dynamically changing graph $G = (V, E)$:
 - $\text{insert}(u, v)$: inserts edge (u, v) in E
 - $\text{delete}(u, v)$: deletes edge (u, v) from E
 - $\text{connected}(u, v)$: reports whether vertices u and v are connected in G
- We refer to insert and delete as update operations and to connected as a query operation
- Initial graph: $|V| = n$ vertices, $E = \emptyset$
- Updates and queries are revealed one by one in an online sequence
- We give a data structure with:
 - $O(\log n)$ worst-case query time
 - $O(\log^2 n)$ amortized update time

Edge Levels and Clusters

- Our data structure will maintain a *level* $\ell(e)$ for each $e \in E$ where $0 \leq \ell(e) \leq \ell_{\max} = \lfloor \log n \rfloor$
- For $0 \leq i \leq \ell_{\max}$, let $G_i = (V, E_i)$ denote the subgraph of G containing edges e with $\ell(e) \geq i$
- We have $E = E_0 \supseteq E_1 \supseteq E_2 \supseteq \dots \supseteq E_{\ell_{\max}}$
- Connected components of G_i are called *i-clusters* or just *clusters*
- Invariant: any *i-cluster* contains at most $\lfloor n/2^i \rfloor$ vertices
- 0-clusters are the connected components of G
- ℓ_{\max} -clusters are vertices of V (why?)


Clusters



Cluster forest


- The *cluster forest* of G is a forest \mathcal{C} of rooted trees where each node u corresponds to a cluster $C(u)$
- A node u at level $i < \ell_{\max}$ has as children the level $(i + 1)$ -nodes v such that $C(v) \subseteq C(u)$
- Roots of \mathcal{C} correspond to connected components of G and leaves of \mathcal{C} correspond to vertices of G
- Each node u of \mathcal{C} is associated with its *size* $n(u)$ which is the number of leaves in the subtree of \mathcal{C} rooted at u

Cluster forest




Answering Queries

- To determine if vertices u and v are connected in G , traverse the leaf-to-root paths from u and v in \mathcal{C}
- Then u and v are connected in G iff the roots are the same
- Query time $O(\log n)$




Handling $\text{insert}(u, v)$

- Initialize $\ell(u, v) \leftarrow 0$
- r_u, r_v : roots of trees of \mathcal{C} containing u and v , respectively
- If $r_u = r_v$, \mathcal{C} is not changed
- Otherwise, r_u and r_v are *merged*
- This corresponds to merging $C(r_u)$ and $C(r_v)$




Handling $\text{insert}(u, v)$

- Initialize $\ell(u, v) \leftarrow 0$
- r_u, r_v : roots of trees of \mathcal{C} containing u and v , respectively
- If $r_u = r_v$, \mathcal{C} is not changed
- Otherwise, r_u and r_v are *merged*
- This corresponds to merging $C(r_u)$ and $C(r_v)$




Handling $\text{insert}(u, v)$

- Initialize $\ell(u, v) \leftarrow 0$
- r_u, r_v : roots of trees of \mathcal{C} containing u and v , respectively
- If $r_u = r_v$, \mathcal{C} is not changed
- Otherwise, r_u and r_v are *merged*
- This corresponds to merging $C(r_u)$ and $C(r_v)$



Handling $\text{insert}(u, v)$


- Initialize $\ell(u, v) \leftarrow 0$
- r_u, r_v : roots of trees of \mathcal{C} containing u and v , respectively
- If $r_u = r_v$, \mathcal{C} is not changed
- Otherwise, r_u and r_v are *merged*
- This corresponds to merging $C(r_u)$ and $C(r_v)$




Handling $\text{delete}(u, v)$

- Let $i = \ell(u, v)$ and let C_u and C_v be the $(i + 1)$ -clusters containing u and v
- Assume $C_u \neq C_v$ since otherwise, \mathcal{C} is not changed
- Let M_i be the multigraph with $(i + 1)$ -clusters as vertices and level i -edges of G as edges
- In M_i , execute two standard search procedures in parallel, one starting in C_u , the other starting in C_v
- Terminate both procedures when in one of the following two cases:
 - a vertex of M_i is explored by both search procedures
 - one of the search procedures has no more edges to explore


A vertex of M_i is explored by both search procedures




A vertex of M_i is explored by both search procedures




A vertex of M_i is explored by both search procedures




A vertex of M_i is explored by both search procedures




A vertex of M_i is explored by both search procedures




A vertex of M_i is explored by both search procedures




A vertex of M_i is explored by both search procedures




A vertex of M_i is explored by both search procedures




A search procedure has no more edges to explore




A search procedure has no more edges to explore




A search procedure has no more edges to explore




A search procedure has no more edges to explore




A search procedure has no more edges to explore



A search procedure has no more edges to explore




A search procedure has no more edges to explore




Updates to \mathcal{C}

- If the two search procedures meet, the level i -cluster C containing (u, v) is still connected so C remains a level i -cluster




Updates to \mathcal{C}

- If the two search procedures meet, the level i -cluster C containing (u, v) is still connected so C remains a level i -cluster
- Otherwise, C is split in two, one part containing C_u and one containing C_v



Updates to \mathcal{C}


- If the two search procedures meet, the level i -cluster C containing (u, v) is still connected so C remains a level i -cluster
- Otherwise, C is split in two, one part containing C_u and one containing C_v



- If $i > 0$, recurse on level $i - 1$


Edge Level Increases after Search

- Recall: each node w of \mathcal{C} is associated with its size $n(w)$
- $n(w)$ is the number of vertices of V in cluster $C(w)$
- For the search procedure that explored clusters of smallest total size, all its visited edges have their levels increased




Edge Level Increases after Search

- Recall: each node w of \mathcal{C} is associated with its size $n(w)$
- $n(w)$ is the number of vertices of V in cluster $C(w)$
- For the search procedure that explored clusters of smallest total size, all its visited edges have their levels increased




Edge Level Increases after Search

- Recall: each node w of \mathcal{C} is associated with its size $n(w)$
- $n(w)$ is the number of vertices of V in cluster $C(w)$
- For the search procedure that explored clusters of smallest total size, all its visited edges have their levels increased




Edge Level Increases after Search

- Recall: each node w of \mathcal{C} is associated with its size $n(w)$
- $n(w)$ is the number of vertices of V in cluster $C(w)$
- For the search procedure that explored clusters of smallest total size, all its visited edges have their levels increased




Edge Level Increases after Search

- Recall: each node w of \mathcal{C} is associated with its size $n(w)$
- $n(w)$ is the number of vertices of V in cluster $C(w)$
- For the search procedure that explored clusters of smallest total size, all its visited edges have their levels increased




Edge Level Increases after Search

- Recall: each node w of \mathcal{C} is associated with its size $n(w)$
- $n(w)$ is the number of vertices of V in cluster $C(w)$
- For the search procedure that explored clusters of smallest total size, all its visited edges have their levels increased



Maintaining the Invariant

- Parent level i -cluster C has size at most $\lfloor n/2^i \rfloor$




- The smaller side has size at most $\lfloor n/2^{i+1} \rfloor$ since otherwise, C would have size $\geq 2\lceil n/2^{i+1} \rceil > 2 \cdot n/2^{i+1} \geq \lfloor n/2^i \rfloor$
- Thus, the invariant is still satisfied after merging level $(i + 1)$ -clusters

Overall Amortized Analysis

- Suppose each search procedure uses $O(1)$ time per edge visited
- For the analysis, we let each edge pay $O(1)$ credits when its level is increased
- The search on the smaller side is thus paid for by its visited edges
- The other search visits the same number of edges (plus/minus 1)
- Hence, the edge level increases can pay for both search procedures
- Max level of an edge: $\ell_{\max} = \lfloor \log n \rfloor = O(\log n)$
- Amortized time per update is thus $O(\log n)$
- What is the problem with this analysis?
 - The multigraph M_i is not stored explicitly
 - Thus, we cannot ensure $O(1)$ time per edge visited
 - We will instead show how to get $O(\log n)$ time per edge visited
 - This will give $O(\log^2 n)$ amortized update time


Traversing a single graph edge

Tree in cluster forest \mathcal{C}



Assuming a Binary Cluster Forest \mathcal{C}

- Assume \mathcal{C} is binary: every node has at most two children
- At each such node u , store an ℓ_{\max} -bit word, $\text{edge}(u)$
- The i th bit $\text{edge}(u)[i]$ is 1 if and only if a level i -edge of E is incident to a leaf of the subtree of \mathcal{C} rooted at u
- Example with $\ell_{\max} = 5$:



Assuming a Binary Cluster Forest \mathcal{C}


- Assume \mathcal{C} is binary: every node has at most two children
- At each such node u , store an ℓ_{\max} -bit word, $\text{edge}(u)$
- The i th bit $\text{edge}(u)[i]$ is 1 if and only if a level i -edge of E is incident to a leaf of the subtree of \mathcal{C} rooted at u
- Maintaining these bitmaps can be done efficiently (exercise)
- Since \mathcal{C} is binary, we can traverse a single edge of a multigraph in $O(\log n)$ time using the edge-bit maps (how?)
- This gives the desired time bound for the search procedures
- However, we need to deal with the case where \mathcal{C} is not binary

Node Ranks

- Recall: for each node u in \mathcal{C} , $n(u)$ is the number of leaves in the subtree of \mathcal{C} rooted at u
- Define the *rank* of u as $\text{rank}(u) = \lfloor \lg n(u) \rfloor$

Rank Trees

- Let u be a non-leaf node in \mathcal{C}
- Initialize node set R as the children of u in \mathcal{C}
- *Rank trees* of u are formed by repeating the following procedure as long as two nodes of R have the same rank:
 - Remove from R two nodes r_1 and r_2 with $\text{rank}(r_1) = \text{rank}(r_2)$
 - Attach r_1 and r_2 to a parent r of rank $\text{rank}(r) = \text{rank}(r_1) + 1$
 - Add r to R




Local trees

- Let r_1, r_2, \dots, r_k be the final set of rank tree roots in R ordered by decreasing rank:

$$\text{rank}(r_1) > \text{rank}(r_2) > \dots > \text{rank}(r_k)$$

- Local tree $L(u)$ for $k = 5$:*



- Replace edges from u to its children in \mathcal{C} by $L(u)$
- Doing this for all u turns \mathcal{C} into forest \mathcal{C}_L of binary trees

Properties of \mathcal{C}_L

- \mathcal{C}_L has height $O(\log n)$ (exercise)
- Merging nodes u and v in \mathcal{C} involves merging $L(u)$ and $L(v)$ in \mathcal{C}_L
- Splitting a node u involves splitting $L(u)$
- This can be done in $O(\log n)$ time per merge/split and will not increase the asymptotic update time (exercise)

Performance of data structure

- Each edge pays $O(\log n)$ credits each time its level increases
- Its level can never decrease
- Number of levels: $O(\log n)$
- Amortized time per update: $O(\log^2 n)$
- Query time: $O(\log n)$
- Space: $O(m + n \log n)$ words
- Can be improved to $O(m + n)$ by compressing paths in \mathcal{C}_L , whose interior nodes have degree 2, to single edges
- Using a more complicated data structure, both update and query time can be improved by a factor of $\log \log n$
- This is still the fastest deterministic data structure known